

# Report

Date of Submission: 14-11-2024

**Entry Number:** 2024AIB2286, 2024AIB2287, 2024CSY7560, 2024CSY7558

## Objective Function

The objective is to give only those predictions which are with high confidence and developing machine learning models that achieve **high overall accuracy** while ensuring **high-confidence predictions** using the CIFAR-100 dataset that was provided as part of the assignment.

## Model Architecture:

We have tried multiple models using pytorch train and validation data so as to find test accuracy and check which model performs better. The model performing better in pytorch data was then later used for training on kaggle dataset.

Models used are as below:

efficientnet-B0, efficientnet-B3, efficientnet-Lite, Enhanced CNN (custom CNN model), resnet-18, wide-resnet and other custom CNN models.

The epoch wise loss and test accuracy of models can be found in the later part of the report. It was observed that wide-resnet was converging to the better accuracy faster and as training data was small and 5 hours of training time so we went ahead with wide-resnet model

## Approach used while training the model

Hyper parameter tuning (where parameters like batch size, learning rate and early stopping) was changed and accordingly accuracy was detected.

We tried Adam and multiple other optimizers and found that Adam is going ahead.

We tried schedulers like stepLR, cyclicLR, cosine annealing and so found that cyclicLR was performing better and hence went ahead with cyclicLR

We tried multiple loss functions from crossEntropyLoss, focalLoss and even tried making a custom loss function which gives 0 loss for correct classification and -5 loss for false prediction. Found that crossEntropyLoss outputs are better and give more prediction scores.

As the training dataset is small, We tried k-fold cross validation on the kaggle platform but it was taking too much time to train. As we had limited GPU hours available on kaggle, we decided not to use k-fold cross validation.

We also tried multiple data transformation like transforms.RandomCrop(32, padding=4), transforms.RandomHorizontalFlip(), transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.2), transforms.Normalize((0.5071, 0.4865, 0.4409), (0.2673, 0.2564, 0.2762)) etc. along with there permutation and combination.

We went ahead with random crop and random horizontal flip only as other transforms were making predictions tough.

We found that more data would be better as the model has reached its max and requires some more data or other data transformation so as to get better accuracy so as to do so, we added data augmentation transformation which increased the accuracy of model from 75% to 77%.

As the earlier 2 submissions were with penalty term as 0.7 and later was with 0.99. In Order to predict only when confidence is best, we also added temperature scaling and then threshold before prediction of the class.

## Justification of Approach and Experiments:

As explained above We have tried different variations but decided to go with the architecture which was giving best results after training and optimising the same. We have attached the various variations tried and results. One of such experiment is given as below:

### WRN-28-10 Training on CIFAR-100

We tested it out with two different loss functions – Cross-Entropy and Focal Loss – just to see what would happen. We used a WideResNet-28-10,. It's got 28 layers and 10 times wider than your usual ResNet . This architecture is known for doing pretty well with image classification tasks.

- **Training Details:** We trained for 150 epochs with a batch size of 128. We used the Adam optimizer with a learning rate of 0.01 and a scheduler to decrease the learning rate over time.

**What We Found:**

- **Training Loss:** Both loss functions did a good job of minimising the training loss. It dropped significantly in the beginning and then gradually levelled off. Cross-entropy reached a loss of 0.00021, while Focal Loss got down to 0.000069.
- **Time:** Each training run took about 5-6 hours to complete.

#### Things to Improve:

- To check validation accuracy
- **Try different parameters**
- **Hyperparameter Tuning**

- As the training dataset is small, We tried k-fold cross validation on kaggle platform but it was taking too much time to train. As we had limited GPU hours available on kaggle, we decided not to use k-fold cross validation
- We were not able to use various optimization techniques due to resource constraints such as GridSearch, Random Search or PSO. But the parameters have been found after multiple training sessions on different kaggle notebooks.

## Use of Open Source Libraries, codes and Models

- Libraries and Tools: Torch, PIL, NumPy
- Weight Initialisations for models
- Additional Resources

## Training

- **Data Preparation**  
The pickle file
- **Training Methodology**

## Inference

- **Prediction Methodology:**  
If Predicted probability greater or equal to than 0.99 give the prediction else -1
- **Confidence Threshold** : 0.99

## Calibration techniques:

As the training time is too much and due to lack of GPU resources. We could not try GridSearch, RandomSearch or Particle Swarm Optimization techniques for hyper parameters.

Instead all 4 team members used their free quota on kaggle platform to train the models on different parameters and compared and combined the results. Our Submission score started from -28000 and the latest kaggle score is more than 6400.

We have decided to write each finding in excel and communicate with a common whatsapp group.

Variations tried:

Type	Variations Tried
Model	efficientnet-B0, efficientnet-B3, efficientnet-Lite, Enhanced CNN (custom CNN model), resnet-18, VGG-16, VGG-19, MobileNetV2, wide-resnet, resnext and other custom CNN models.
Optimizer	Adam, SGD
Loss Function	Cross Entropy, Penalty for wrong classification
Transforms	RandomCrop, HorizontalFlip, ColorJitter, RandomErasing,
Schedulers	StepLR, CyclicLR, CosineAnnealingLR

## Pre processing Steps

Using the torch library transform sub module, we have applied RandomCrop and RandomHorizontalFlip as the transform to the training data but the same has not been applied to the test model. We also added normalisation and data augmentation so as to increase accuracy and help models to learn features.

## Post-processing Steps

Test model has not been applied to any transform apart from normalisation, which was applied when training.

## Custom CNN model:

```
class DeeperEnhancedCNN(nn.Module):
    def __init__(self):
        super(DeeperEnhancedCNN, self).__init__()
        # First block
        self.conv1 = nn.Conv2d(3, 128, kernel_size=3, stride=1, padding=1)
        self.bn1 = nn.BatchNorm2d(128)
        self.conv2 = nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1)
        self.bn2 = nn.BatchNorm2d(256)
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)

        # Second block
        self.conv3 = nn.Conv2d(256, 512, kernel_size=3, stride=1, padding=1)
        self.bn3 = nn.BatchNorm2d(512)
        self.conv4 = nn.Conv2d(512, 1024, kernel_size=3, stride=1, padding=1)
        self.bn4 = nn.BatchNorm2d(1024)
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)

        # Third block - Added more layers here
        self.conv5 = nn.Conv2d(1024, 2048, kernel_size=3, stride=1, padding=1)
        self.bn5 = nn.BatchNorm2d(2048)
        self.conv6 = nn.Conv2d(2048, 4096, kernel_size=3, stride=1, padding=1) # Extra
convolution
        self.bn6 = nn.BatchNorm2d(4096)
        self.pool3 = nn.MaxPool2d(kernel_size=2, stride=2)

        self.relu = nn.PReLU()
        self.dropout = nn.Dropout(p=0.4)

        # Adjusted the fully connected layer
        self.fc1 = nn.Linear(4096 * 4 * 4, 4096)
        self.fc2 = nn.Linear(4096, 1024)
        self.fc3 = nn.Linear(1024, 100) # For CIFAR-100

    def forward(self, x):
        # First block
        x = self.pool1(self.relu(self.bn1(self.conv1(x))))
        x = self.relu(self.bn2(self.conv2(x)))

        # Second block
        x = self.pool2(self.relu(self.bn3(self.conv3(x))))
        x = self.relu(self.bn4(self.conv4(x)))

        # Third block
```

```
x = self.pool3(self.relu(self.bn5(self.conv5(x))))
x = self.relu(self.bn6(self.conv6(x)))
```

```
# Flatten the tensor
x = x.view(x.size(0), -1)
```

```
# Fully connected layers
x = self.relu(self.fc1(x))
x = self.dropout(x)
x = self.relu(self.fc2(x))
x = self.dropout(x)
x = self.fc3(x)
```

```
return x
```

The above model was ran with batch size 128 till 100 epochs and max accuracy with test set was found to be 57%.

## **Wideresnet:**

# Define the Basic Block used in WideResNet

```
class BasicBlock(nn.Module):
```

```
    def __init__(self, in_planes, out_planes, stride, drop_rate=0.3):
```

```
        super(BasicBlock, self).__init__()
```

```
        self.bn1 = nn.BatchNorm2d(in_planes)
```

```
        self.relu = nn.ReLU(inplace=True)
```

```
        self.conv1 = nn.Conv2d(in_planes, out_planes, kernel_size=3, stride=stride,
padding=1, bias=False)
```

```
        self.bn2 = nn.BatchNorm2d(out_planes)
```

```
        self.conv2 = nn.Conv2d(out_planes, out_planes, kernel_size=3, stride=1, padding=1,
bias=False)
```

```
        self.drop_rate = drop_rate
```

```
        self.equalInOut = (in_planes == out_planes)
```

```
        self.convShortcut = (not self.equalInOut) and nn.Conv2d(in_planes, out_planes,
kernel_size=1, stride=stride, padding=0, bias=False) or None
```

```
    def forward(self, x):
```

```
        out = self.relu(self.bn1(x))
```

```
        if not self.equalInOut:
```

```
            x = out
```

```
        out = self.relu(self.bn2(self.conv1(out)))
```

```
        if self.drop_rate > 0:
```

```
            out = F.dropout(out, p=self.drop_rate, training=self.training)
```

```

out = self.conv2(out)
return torch.add(x if self.convShortcut is None else self.convShortcut(x), out)

```

# Define Network Block

```

class NetworkBlock(nn.Module):
    def __init__(self, nb_layers, in_planes, out_planes, block, stride, drop_rate=0.3):
        super(NetworkBlock, self).__init__()
        self.layer = self._make_layer(block, in_planes, out_planes, nb_layers, stride,
drop_rate)

    def _make_layer(self, block, in_planes, out_planes, nb_layers, stride, drop_rate):
        layers = []
        for i in range(nb_layers):
            layers.append(block(i == 0 and in_planes or out_planes, out_planes, i == 0 and
stride or 1, drop_rate))
        return nn.Sequential(*layers)

    def forward(self, x):
        return self.layer(x)

```

# Define the WideResNet Model

```

class WideResNet(nn.Module):
    def __init__(self, depth, num_classes, widen_factor=1, drop_rate=0.3):
        super(WideResNet, self).__init__()
        self.in_planes = 16
        assert (depth - 4) % 6 == 0, 'Depth should be 6n+4'
        n = (depth - 4) // 6
        k = widen_factor
        nStages = [16, 16 * k, 32 * k, 64 * k]

        self.conv1 = nn.Conv2d(3, nStages[0], kernel_size=3, stride=1, padding=1,
bias=False)
        self.block1 = NetworkBlock(n, nStages[0], nStages[1], BasicBlock, 1, drop_rate)
        self.block2 = NetworkBlock(n, nStages[1], nStages[2], BasicBlock, 2, drop_rate)
        self.block3 = NetworkBlock(n, nStages[2], nStages[3], BasicBlock, 2, drop_rate)
        self.bn1 = nn.BatchNorm2d(nStages[3])
        self.relu = nn.ReLU(inplace=True)
        self.fc = nn.Linear(nStages[3], num_classes)
        self.drop_rate = drop_rate

        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='relu')

    def forward(self, x):
        out = self.conv1(x)
        out = self.block1(out)
        out = self.block2(out)
        out = self.block3(out)

```

```

out = self.block3(out)
out = self.relu(self.bn1(out))
out = F.avg_pool2d(out, 8)
out = out.view(-1, self.fc.in_features)
return self.fc(out)

```

The above model was ran with 128 batch size and till convergence. Below is the loss and accuracy with test set observed

wideresnet model:

Epoch 1/500, Loss: 3.8958370655089083, accuracy: 13.39% Epoch 2/500, Loss: 3.1907375718626523, accuracy: 23.90% Epoch 3/500, Loss: 2.66034812207722, accuracy: 33.61% Epoch 4/500, Loss: 2.2760456333989683, accuracy: 39.56% Epoch 5/500, Loss: 1.980759484078878, accuracy: 42.98% Epoch 6/500, Loss: 1.75116037956589, accuracy: 46.99% Epoch 7/500, Loss: 1.5703523268785013, accuracy: 51.47% Epoch 8/500, Loss: 1.416255613727033, accuracy: 53.48% Epoch 9/500, Loss: 1.2821660611940466, accuracy: 56.87% Epoch 10/500, Loss: 1.1641356271246206, accuracy: 56.43% Epoch 11/500, Loss: 1.053151973982906, accuracy: 57.16% Epoch 12/500, Loss: 0.9569692506509668, accuracy: 58.58% Epoch 13/500, Loss: 0.8613696657788114, accuracy: 60.20% Epoch 14/500, Loss: 0.7829350902296394, accuracy: 59.80% Epoch 15/500, Loss: 0.7034277122496339, accuracy: 60.35% Epoch 16/500, Loss: 0.6365693751198557, accuracy: 61.36% Epoch 17/500, Loss: 0.5702269039190638, accuracy: 62.70% Epoch 18/500, Loss: 0.5269172452294918, accuracy: 59.15% Epoch 19/500, Loss: 0.47140407493657166, accuracy: 61.18% Epoch 20/500, Loss: 0.42413251654571277, accuracy: 62.86% Epoch 21/500, Loss: 0.3901892089858994, accuracy: 61.83% Epoch 22/500, Loss: 0.3585834849597243, accuracy: 60.09% Epoch 23/500, Loss: 0.3293017967963767, accuracy: 62.20%

## **Resnet 18:**

```
model = models.resnet18(pretrained=False)
```

The model output where observed as below:

Epoch 1/100, Loss: 4.194681015160993  
Validation Accuracy: 11.22%  
Epoch 2/100, Loss: 3.710814584551565  
Validation Accuracy: 16.08%  
Epoch 3/100, Loss: 3.4643728909894937  
Validation Accuracy: 18.89%  
Epoch 4/100, Loss: 3.2846149745804576  
Validation Accuracy: 21.70%



Epoch 5/100, Loss: 3.113571018209238  
Validation Accuracy: 23.64%  
Epoch 6/100, Loss: 2.9739636545595913  
Validation Accuracy: 26.10%  
Epoch 7/100, Loss: 2.8548560099833455  
Validation Accuracy: 27.69%  
Epoch 8/100, Loss: 2.743235691124216  
Validation Accuracy: 30.38%  
Epoch 9/100, Loss: 2.647169861342291  
Validation Accuracy: 31.57%  
Epoch 10/100, Loss: 2.569054919435545  
Validation Accuracy: 32.60%  
Epoch 11/100, Loss: 2.4967064083079853  
Validation Accuracy: 33.25%  
Epoch 12/100, Loss: 2.4184513579853966  
Validation Accuracy: 34.98%  
Epoch 13/100, Loss: 2.3558037220059758  
Validation Accuracy: 35.83%  
Epoch 14/100, Loss: 2.2940699980996757  
Validation Accuracy: 36.07%  
Epoch 15/100, Loss: 2.2353271904503904  
Validation Accuracy: 37.40%  
Epoch 16/100, Loss: 2.1800304642113884  
Validation Accuracy: 37.41%  
Epoch 17/100, Loss: 2.1265135479095343  
Validation Accuracy: 38.73%  
Epoch 18/100, Loss: 2.0913470681670985  
Validation Accuracy: 38.61%  
Epoch 19/100, Loss: 2.03326692666544  
Validation Accuracy: 38.92%  
Epoch 20/100, Loss: 1.9872566513393237  
Validation Accuracy: 39.70%  
Epoch 21/100, Loss: 1.9433395149152908  
Validation Accuracy: 40.77%  
Epoch 22/100, Loss: 1.8978181278614132  
Validation Accuracy: 41.37%  
Epoch 23/100, Loss: 1.8598604290686604  
Validation Accuracy: 41.42%  
Epoch 24/100, Loss: 1.813386885101533  
Validation Accuracy: 41.48%  
Epoch 25/100, Loss: 1.772424898793935  
Validation Accuracy: 41.60%  
Epoch 26/100, Loss: 1.7409106005183266  
Validation Accuracy: 42.30%  
Epoch 27/100, Loss: 1.6968459147016715  
Validation Accuracy: 43.67%  
Epoch 28/100, Loss: 1.6643307340114624  
Validation Accuracy: 43.05%

Epoch 29/100, Loss: 1.6227991324861337  
Validation Accuracy: 42.57%  
Epoch 30/100, Loss: 1.589911506913812  
Validation Accuracy: 43.39%  
Early stopping triggered.  
Test Accuracy: 42.56%  
Execution time: 959.2236661911011

## **Efficientnetlite:**

```
class EfficientNetLiteCNN(nn.Module):
    def __init__(self):
        super(EfficientNetLiteCNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1, stride=1)
        self.bn1 = nn.BatchNorm2d(32)
        self.dwconv2 = nn.Conv2d(32, 32, kernel_size=3, padding=1, groups=32)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=1)
        self.bn2 = nn.BatchNorm2d(64)
        self.dwconv3 = nn.Conv2d(64, 64, kernel_size=3, padding=1, groups=64)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=1)
        self.bn3 = nn.BatchNorm2d(128)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(128 * 8 * 8, 512)
        self.dropout = nn.Dropout(0.5)
        self.fc2 = nn.Linear(512, 100)

    def forward(self, x):
        x = torch.relu(self.bn1(self.conv1(x)))
        x = torch.relu(self.bn2(self.conv2(self.dwconv2(x))))
        x = self.pool(x)
        x = torch.relu(self.bn3(self.conv3(self.dwconv3(x))))
        x = self.pool(x)
        x = x.view(-1, 128 * 8 * 8)
        x = torch.relu(self.fc1(x))
        x = self.dropout(x)
        x = self.fc2(x)
        return x

# Hyperparameters
batch_size = 128
learning_rate = 0.001
```

Below are the observation found:

efficient net lite model:

Epoch 1/500, Loss: 4.360215240732178, accuracy: 10.19%Epoch 2/500, Loss: 4.027833084925971, accuracy: 13.71%Epoch 3/500, Loss: 3.8574972786866795, accuracy: 15.93%Epoch 4/500, Loss: 3.739587589907829, accuracy: 17.53%Epoch 5/500, Loss: 3.6378826143796488, accuracy: 19.27%Epoch 6/500, Loss: 3.5537267566641884, accuracy: 21.19%Epoch 7/500, Loss: 3.476924180374731, accuracy: 21.70%Epoch 8/500, Loss: 3.408281158608244, accuracy: 23.11%Epoch 9/500, Loss: 3.340460462033596, accuracy: 24.61%Epoch 10/500, Loss: 3.286372683542159, accuracy: 25.66%Epoch 11/500, Loss: 3.228576555276466, accuracy: 26.80%Epoch 12/500, Loss: 3.172423386512815, accuracy: 27.75%Epoch 13/500, Loss: 3.138412274363096, accuracy: 28.80%Epoch 14/500, Loss: 3.1039964268579507, accuracy: 28.94%Epoch 15/500, Loss: 3.065286183296262, accuracy: 29.72%Epoch 16/500, Loss: 3.0393613207980494, accuracy: 30.62%Epoch 17/500, Loss: 3.015606296031981, accuracy: 30.00%Epoch 18/500, Loss: 2.986381257586467, accuracy: 30.73%Epoch 19/500, Loss: 2.968880922288236, accuracy: 31.91%Epoch 20/500, Loss: 2.946484629760313, accuracy: 32.48%Epoch 21/500, Loss: 2.9324763790725745, accuracy: 32.38%Epoch 22/500, Loss: 2.91599942229288, accuracy: 32.97%Epoch 23/500, Loss: 2.891361853655647, accuracy: 33.38%Epoch 24/500, Loss: 2.878907850636241, accuracy: 32.76%Epoch 25/500, Loss: 2.862930344803559, accuracy: 34.01%Epoch 26/500, Loss: 2.8489256868581943, accuracy: 33.72%Epoch 27/500, Loss: 2.836324024078486, accuracy: 34.62%Epoch 28/500, Loss: 2.827077588766737, accuracy: 35.48%Epoch 29/500, Loss: 2.800273266594733, accuracy: 35.27%Epoch 30/500, Loss: 2.799769186302829, accuracy: 35.07%Epoch 31/500, Loss: 2.7807259108404367, accuracy: 35.97%Epoch 32/500, Loss: 2.7779379732468548, accuracy: 35.48%Epoch 33/500, Loss: 2.7674035788192164, accuracy: 36.27%Epoch 34/500, Loss: 2.7540074364303626, accuracy: 37.04%Epoch 35/500, Loss: 2.7438957666801977, accuracy: 36.99%Epoch 36/500, Loss: 2.726191955454209, accuracy: 36.47%Epoch 37/500, Loss: 2.729940025397884, accuracy: 36.88%Epoch 38/500, Loss: 2.7138825700715983, accuracy: 37.44%Epoch 39/500, Loss: 2.710140200832006, accuracy: 36.63%Epoch 40/500, Loss: 2.694559768642611, accuracy: 37.98%Epoch 41/500, Loss: 2.6931150307130935, accuracy: 38.27%Epoch 42/500, Loss: 2.6897601581290553, accuracy: 37.46%Epoch 43/500, Loss: 2.6791635638917497, accuracy: 38.33%Epoch 44/500, Loss: 2.657219436772339, accuracy: 38.35%Epoch 45/500, Loss: 2.6526657004490533, accuracy: 38.09%Epoch 46/500, Loss: 2.6493101705370656, accuracy: 38.96%Epoch 47/500, Loss: 2.6280903041820087, accuracy: 38.76%Epoch 48/500, Loss: 2.628377166855366, accuracy: 39.72%Epoch 49/500, Loss: 2.6115937623221552, accuracy: 38.53%Epoch 50/500, Loss: 2.614183448464669, accuracy: 40.15%Epoch 51/500, Loss: 2.6038029352417382, accuracy: 38.83%Epoch 52/500, Loss: 2.6081756447892053, accuracy: 40.55%Epoch 53/500, Loss: 2.5963790099639112, accuracy: 39.75%Epoch 54/500, Loss: 2.582945385247545, accuracy: 40.82%Epoch 55/500, Loss: 2.579683523348835, accuracy: 40.78%Epoch 56/500, Loss: 2.5666697872874074, accuracy: 40.06%Epoch 57/500, Loss: 2.5633385998513694, accuracy: 39.17%Epoch 58/500, Loss: 2.559815848270036, accuracy: 41.18%Epoch 59/500, Loss: 2.5505428295916, accuracy: 40.18%Epoch 60/500, Loss: 2.540532194440017, accuracy: 40.55%Epoch 61/500, Loss: 2.528377365883049, accuracy: 41.00%Epoch 62/500, Loss: 2.527028257584633, accuracy: 40.89%Epoch 63/500, Loss: 2.5111712324040014, accuracy: 41.87%Epoch 64/500, Loss: 2.515386178365449, accuracy: 41.88%Epoch 65/500, Loss: 2.5021812751165133, accuracy: 41.49%Epoch 66/500, Loss: 2.493788082581347,

accuracy: 41.70%Epoch 67/500, Loss: 2.5010043611306974, accuracy: 41.47%Epoch 68/500, Loss: 2.4815885093815795, accuracy: 41.64%Epoch 69/500, Loss: 2.475962230311635, accuracy: 42.38%Epoch 70/500, Loss: 2.471860908181466, accuracy: 41.34%Epoch 71/500, Loss: 2.448628339011346, accuracy: 41.62%Epoch 72/500, Loss: 2.4667965438969603, accuracy: 40.74%Epoch 73/500, Loss: 2.462963899383155, accuracy: 42.62%Epoch 74/500, Loss: 2.4485643702699704, accuracy: 42.27%Epoch 75/500, Loss: 2.437971852929391, accuracy: 42.93%Epoch 76/500, Loss: 2.4348993405051855, accuracy: 43.14%Epoch 77/500, Loss: 2.4318822302171945, accuracy: 42.63%Epoch 78/500, Loss: 2.428695847311288, accuracy: 42.56%Epoch 79/500, Loss: 2.416344038970635, accuracy: 42.63%Epoch 80/500, Loss: 2.4110349538686027, accuracy: 42.75%Epoch 81/500, Loss: 2.4078052007328825, accuracy: 43.66%Epoch 82/500, Loss: 2.392879025710513, accuracy: 42.95%Epoch 83/500, Loss: 2.3909472004531898, accuracy: 43.54%Epoch 84/500, Loss: 2.3889439816365154, accuracy: 42.57%Epoch 85/500, Loss: 2.3828033955810626, accuracy: 43.94%Epoch 86/500, Loss: 2.3701141920236064, accuracy: 43.97%Epoch 87/500, Loss: 2.372822029511337, accuracy: 43.98%Epoch 88/500, Loss: 2.3637990914952116, accuracy: 43.31%Epoch 89/500, Loss: 2.3598635376566817, accuracy: 43.09%Epoch 90/500, Loss: 2.3472682024206954, accuracy: 44.15%Epoch 91/500, Loss: 2.360951248947007, accuracy: 43.62%Epoch 92/500, Loss: 2.3421943681624233, accuracy: 43.96%Epoch 93/500, Loss: 2.329005687742892, accuracy: 44.26%Epoch 94/500, Loss: 2.332755136367915, accuracy: 44.56%Epoch 95/500, Loss: 2.3276535437235135, accuracy: 45.47%Epoch 96/500, Loss: 2.319227936932498, accuracy: 44.42%Epoch 97/500, Loss: 2.31093358078881, accuracy: 45.11%Epoch 98/500, Loss: 2.3100625118026343, accuracy: 44.93%Epoch 99/500, Loss: 2.313205302218952, accuracy: 45.06%Epoch 100/500, Loss: 2.309698928652517, accuracy: 44.90%Epoch 101/500, Loss: 2.2937696568496393, accuracy: 44.80%Epoch 102/500, Loss: 2.2878667171044116, accuracy: 45.31%Epoch 103/500, Loss: 2.2849646065851004, accuracy: 44.64%Epoch 104/500, Loss: 2.279021483248152, accuracy: 44.89%Epoch 105/500, Loss: 2.2824555348862163, accuracy: 43.38%

## **MobileNet - v2:**

model = models.mobilenet\_v2(pretrained=False)

Below are the observation found for model:

Epoch 1/500, Loss: 4.603856518445418, accuracy: 2.09%Epoch 2/500, Loss: 4.514468612573336, accuracy: 2.59%Epoch 3/500, Loss: 4.448392982678035, accuracy: 3.54%Epoch 4/500, Loss: 4.381721428288218, accuracy: 3.92%Epoch 5/500, Loss: 4.297346371214103, accuracy: 5.24%Epoch 6/500, Loss: 4.217047803542194, accuracy: 6.52%Epoch 7/500, Loss: 4.138686625244063, accuracy: 7.18%Epoch 8/500, Loss: 4.067462747968981, accuracy: 7.89%Epoch 9/500, Loss: 4.007351367979708, accuracy: 9.16%Epoch 10/500, Loss: 3.9429921195330215, accuracy: 9.99%Epoch 11/500, Loss: 3.8775348120638173, accuracy: 10.70%Epoch 12/500, Loss: 3.816184785359961, accuracy: 12.39%Epoch 13/500, Loss: 3.76094287252792, accuracy: 12.99%Epoch 14/500, Loss: 3.6912881443872476, accuracy: 14.29%Epoch 15/500, Loss: 3.634215004913642, accuracy: 15.19%Epoch 16/500, Loss: 3.5748970069543784, accuracy: 15.80%Epoch

17/500, Loss: 3.522492667293305, accuracy: 16.89%Epoch 18/500, Loss: 3.4652341271910214, accuracy: 18.09%Epoch 19/500, Loss: 3.406921383669919, accuracy: 18.35%Epoch 20/500, Loss: 3.361202163769461, accuracy: 19.79%Epoch 21/500, Loss: 3.3026282153166164, accuracy: 20.46%Epoch 22/500, Loss: 3.265655853559294, accuracy: 20.59%Epoch 23/500, Loss: 3.2140242779041497, accuracy: 21.65%Epoch 24/500, Loss: 3.1664164480955703, accuracy: 22.41%Epoch 25/500, Loss: 3.1225082666977593, accuracy: 23.15%Epoch 26/500, Loss: 3.0812546566624164, accuracy: 23.88%Epoch 27/500, Loss: 3.039326575101184, accuracy: 24.92%Epoch 28/500, Loss: 3.001599287437966, accuracy: 24.49%Epoch 29/500, Loss: 2.9614796918981217, accuracy: 25.05%Epoch 30/500, Loss: 2.91814642306179, accuracy: 26.19%Epoch 31/500, Loss: 2.8867737850569704, accuracy: 26.49%Epoch 32/500, Loss: 2.8578186968098516, accuracy: 27.29%Epoch 33/500, Loss: 2.814201757426152, accuracy: 26.95%Epoch 34/500, Loss: 2.7827519455834118, accuracy: 28.35%Epoch 35/500, Loss: 2.751336758703832, accuracy: 28.32%Epoch 36/500, Loss: 2.7173804120944283, accuracy: 29.62%Epoch 37/500, Loss: 2.6782970983354026, accuracy: 29.43%Epoch 38/500, Loss: 2.6444375984504096, accuracy: 30.63%Epoch 39/500, Loss: 2.612551815979316, accuracy: 30.03%Epoch 40/500, Loss: 2.575558686195432, accuracy: 31.12%Epoch 41/500, Loss: 2.544057719237969, accuracy: 32.37%Epoch 42/500, Loss: 2.5215466906652426, accuracy: 31.63%Epoch 43/500, Loss: 2.4730424850493136, accuracy: 32.76%Epoch 44/500, Loss: 2.4589860366128593, accuracy: 32.54%Epoch 45/500, Loss: 2.4271067194926466, accuracy: 34.29%Epoch 46/500, Loss: 2.3912697716442217, accuracy: 33.97%Epoch 47/500, Loss: 2.35498140504598, accuracy: 34.12%Epoch 48/500, Loss: 2.328483898316503, accuracy: 34.36%Epoch 49/500, Loss: 2.304700252650034, accuracy: 34.99%Epoch 50/500, Loss: 2.272432833681326, accuracy: 35.50%Epoch 51/500, Loss: 2.239185177151809, accuracy: 35.47%Epoch 52/500, Loss: 2.21773415971595, accuracy: 35.42%Epoch 53/500, Loss: 2.1925170692946296, accuracy: 36.16%Epoch 54/500, Loss: 2.1667599147542966, accuracy: 36.45%Epoch 55/500, Loss: 2.1440669779887287, accuracy: 37.13%Epoch 56/500, Loss: 2.1143350162164634, accuracy: 37.11%Epoch 57/500, Loss: 2.087872126218303, accuracy: 37.48%Epoch 58/500, Loss: 2.067683346436152, accuracy: 37.27%Epoch 59/500, Loss: 2.038112694040284, accuracy: 37.89%Epoch 60/500, Loss: 2.028090117532579, accuracy: 37.48%Epoch 61/500, Loss: 1.982333795798709, accuracy: 38.56%Epoch 62/500, Loss: 1.9776658213047116, accuracy: 38.03%Epoch 63/500, Loss: 1.9379768063650107, accuracy: 37.43%Epoch 64/500, Loss: 1.920553028431085, accuracy: 39.05%Epoch 65/500, Loss: 1.9018781310152215, accuracy: 38.78%Epoch 66/500, Loss: 1.8866547477214843, accuracy: 38.56%Epoch 67/500, Loss: 1.8526559394338857, accuracy: 39.93%Epoch 68/500, Loss: 1.830819897334594, accuracy: 39.68%Epoch 69/500, Loss: 1.8184470142549871, accuracy: 38.97%Epoch 70/500, Loss: 1.804222505335308, accuracy: 39.74%wide net model:  
Epoch 1/500, Loss: 4.647676211793709, accuracy: 1.49%Epoch 2/500, Loss: 4.585302051680777, accuracy: 2.22%Epoch 3/500, Loss: 4.4861460573533005, accuracy: 3.53%Epoch 4/500, Loss: 4.386408864384722, accuracy: 4.27%Epoch 5/500, Loss: 4.3019383837804765, accuracy: 5.48%Epoch 6/500, Loss: 4.235437116354627, accuracy: 5.74%Epoch 7/500, Loss: 4.152704055961745, accuracy: 6.83%Epoch 8/500, Loss: 4.084733406905933, accuracy: 8.24%Epoch 9/500, Loss: 4.011959139343418, accuracy: 9.46%Epoch 10/500, Loss: 3.9428267594798445, accuracy: 10.50%Epoch 11/500, Loss:

3.8791943767186625, accuracy: 11.37%Epoch 12/500, Loss: 3.8155172999252747, accuracy: 11.90%Epoch 13/500, Loss: 3.7498889134058255, accuracy: 13.78%Epoch 14/500, Loss: 3.6886627454586955, accuracy: 14.12%Epoch 15/500, Loss: 3.6257643724036646, accuracy: 15.09%Epoch 16/500, Loss: 3.5627404149535975, accuracy: 16.13%Epoch 17/500, Loss: 3.5164357018287835, accuracy: 16.70%Epoch 18/500, Loss: 3.454591001086223, accuracy: 17.46%Epoch 19/500, Loss: 3.403981969789471, accuracy: 18.15%Epoch 20/500, Loss: 3.3558685456395456, accuracy: 20.08%Epoch 21/500, Loss: 3.307169325821235, accuracy: 19.82%Epoch 22/500, Loss: 3.254223558908838, accuracy: 21.33%Epoch 23/500, Loss: 3.2058057583811337, accuracy: 21.91%Epoch 24/500, Loss: 3.1551490238560436, accuracy: 22.24%Epoch 25/500, Loss: 3.1076970429676574, accuracy: 23.37%Epoch 26/500, Loss: 3.05836103212498, accuracy: 23.73%Epoch 27/500, Loss: 3.0109008139051743, accuracy: 25.46%Epoch 28/500, Loss: 2.969358307626241, accuracy: 25.02%Epoch 29/500, Loss: 2.9369211202997074, accuracy: 25.93%Epoch 30/500, Loss: 2.8919867963132346, accuracy: 26.40%Epoch 31/500, Loss: 2.853459740538731, accuracy: 27.49%Epoch 32/500, Loss: 2.808312281318333, accuracy: 27.97%Epoch 33/500, Loss: 2.770021930679946, accuracy: 29.48%Epoch 34/500, Loss: 2.730112189222175, accuracy: 29.47%Epoch 35/500, Loss: 2.692394308421923, accuracy: 29.37%Epoch 36/500, Loss: 2.656702952006894, accuracy: 29.71%Epoch 37/500, Loss: 2.621962042415843, accuracy: 30.59%Epoch 38/500, Loss: 2.5940499921588946, accuracy: 31.94%Epoch 39/500, Loss: 2.5587541672884657, accuracy: 32.00%Epoch 40/500, Loss: 2.528397195479449, accuracy: 32.04%Epoch 41/500, Loss: 2.4970059760696137, accuracy: 32.56%Epoch 42/500, Loss: 2.4571884349179083, accuracy: 33.42%Epoch 43/500, Loss: 2.4334481146634386, accuracy: 33.37%Epoch 44/500, Loss: 2.40537991121297, accuracy: 34.35%Epoch 45/500, Loss: 2.3839553472635995, accuracy: 34.58%Epoch 46/500, Loss: 2.3472972480232452, accuracy: 33.78%Epoch 47/500, Loss: 2.3199837909025303, accuracy: 34.57%Epoch 48/500, Loss: 2.29955858068393, accuracy: 35.14%Epoch 49/500, Loss: 2.2767575305441152, accuracy: 35.13%Epoch 50/500, Loss: 2.248543215529693, accuracy: 35.72%Epoch 51/500, Loss: 2.2193917411062722, accuracy: 35.49%Epoch 52/500, Loss: 2.195910219951054, accuracy: 36.07%Epoch 53/500, Loss: 2.1663334205022555, accuracy: 36.85%Epoch 54/500, Loss: 2.149907591702688, accuracy: 36.19%Epoch 55/500, Loss: 2.126648876063354, accuracy: 37.51%Epoch 56/500, Loss: 2.108421998255698, accuracy: 37.44%Epoch 57/500, Loss: 2.080849638985246, accuracy: 37.79%Epoch 58/500, Loss: 2.045118176723685, accuracy: 37.86%Epoch 59/500, Loss: 2.0330254848655835, accuracy: 37.92%Epoch 60/500, Loss: 2.0049837852073145, accuracy: 37.80%Epoch 61/500, Loss: 1.99305193655936, accuracy: 38.33%Epoch 62/500, Loss: 1.9649440870260644, accuracy: 37.84%Epoch 63/500, Loss: 1.9324648063201124, accuracy: 38.62%Epoch 64/500, Loss: 1.9221781644674822, accuracy: 39.20%Epoch 65/500, Loss: 1.9071738314445672, accuracy: 38.96%Epoch 66/500, Loss: 1.8741448565822123, accuracy: 39.23%Epoch 67/500, Loss: 1.8644547151482624, accuracy: 39.30%Epoch 68/500, Loss: 1.8341989419649325, accuracy: 39.04%Epoch 69/500, Loss: 1.810324299670851, accuracy: 39.13%Epoch 70/500, Loss: 1.7878377696742183, accuracy: 38.95%

## **Custom CNN:**

```
model = Sequential()
# 128 and not only 32 filters because there are 100 classes. 32 filters gave bad results.
model.add(Conv2D(128, (3, 3), padding='same', input_shape=X_train.shape[1:]))
model.add(Activation('elu'))

model.add(Conv2D(128, (3, 3)))
model.add(Activation('elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(256, (3, 3), padding='same'))
model.add(Activation('elu'))

model.add(Conv2D(256, (3, 3)))
model.add(Activation('elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(512, (3, 3), padding='same'))
model.add(Activation('elu'))

model.add(Conv2D(512, (3, 3)))
model.add(Activation('elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(1024))
model.add(Activation('elu'))
model.add(Dropout(0.5))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))

model.summary()
```

Below are the observation found:

```
1000/1000 ————— 2s 2ms/step - accuracy:
0.0000e+00 - loss: 0.0000e+00 - val_accuracy: 0.2158 - val_loss: 3.3008
Epoch 3/600
1000/1000 ————— 40s 39ms/step - accuracy: 0.1876
- loss: 3.4240 - val_accuracy: 0.2981 - val_loss: 2.8770
Epoch 4/600
1000/1000 ————— 2s 2ms/step - accuracy:
0.0000e+00 - loss: 0.0000e+00 - val_accuracy: 0.2981 - val_loss: 2.8770
```

Epoch 5/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.2525  
- loss: 3.0732 - val\_accuracy: 0.3439 - val\_loss: 2.6524

Epoch 6/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.3439 - val\_loss: 2.6524

Epoch 7/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.2902  
- loss: 2.8694 - val\_accuracy: 0.3646 - val\_loss: 2.5216

Epoch 8/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.3646 - val\_loss: 2.5216

Epoch 9/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.3290  
- loss: 2.7072 - val\_accuracy: 0.3970 - val\_loss: 2.3511

Epoch 10/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.3970 - val\_loss: 2.3511

Epoch 11/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.3521  
- loss: 2.5788 - val\_accuracy: 0.4164 - val\_loss: 2.2587

Epoch 12/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.4164 - val\_loss: 2.2587

Epoch 13/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.3744  
- loss: 2.4752 - val\_accuracy: 0.4413 - val\_loss: 2.1648

Epoch 14/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.4413 - val\_loss: 2.1648

Epoch 15/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.3833  
- loss: 2.4176 - val\_accuracy: 0.4594 - val\_loss: 2.1087

Epoch 16/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.4594 - val\_loss: 2.1087

Epoch 17/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.4028  
- loss: 2.3433 - val\_accuracy: 0.4738 - val\_loss: 2.0255

Epoch 18/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.4738 - val\_loss: 2.0255

Epoch 19/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4170  
- loss: 2.2627 - val\_accuracy: 0.4769 - val\_loss: 1.9993

Epoch 20/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.4769 - val\_loss: 1.9993



Epoch 21/600  
1000/1000 ————— 40s 40ms/step - accuracy: 0.4273  
- loss: 2.2136 - val\_accuracy: 0.4872 - val\_loss: 1.9420

Epoch 22/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.4872 - val\_loss: 1.9420

Epoch 23/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4411  
- loss: 2.1559 - val\_accuracy: 0.5033 - val\_loss: 1.8873

Epoch 24/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5033 - val\_loss: 1.8873

Epoch 25/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4508  
- loss: 2.1096 - val\_accuracy: 0.5109 - val\_loss: 1.8578

Epoch 26/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5109 - val\_loss: 1.8578

Epoch 27/600  
1000/1000 ————— 40s 40ms/step - accuracy: 0.4566  
- loss: 2.0708 - val\_accuracy: 0.5055 - val\_loss: 1.8593

Epoch 28/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5055 - val\_loss: 1.8593

Epoch 29/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4742  
- loss: 2.0049 - val\_accuracy: 0.5154 - val\_loss: 1.8104

Epoch 30/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5154 - val\_loss: 1.8104

Epoch 31/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4782  
- loss: 1.9715 - val\_accuracy: 0.5338 - val\_loss: 1.7592

Epoch 32/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5338 - val\_loss: 1.7592

Epoch 33/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4887  
- loss: 1.9317 - val\_accuracy: 0.5389 - val\_loss: 1.7319

Epoch 34/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5389 - val\_loss: 1.7319

Epoch 35/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.4923  
- loss: 1.9115 - val\_accuracy: 0.5450 - val\_loss: 1.6767

Epoch 36/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5450 - val\_loss: 1.6767

Epoch 37/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5006  
- loss: 1.8671 - val\_accuracy: 0.5426 - val\_loss: 1.6868

Epoch 38/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5426 - val\_loss: 1.6868

Epoch 39/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5052  
- loss: 1.8473 - val\_accuracy: 0.5504 - val\_loss: 1.6775

Epoch 40/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5504 - val\_loss: 1.6775

Epoch 41/600  
1000/1000 ————— 40s 40ms/step - accuracy: 0.5164  
- loss: 1.7969 - val\_accuracy: 0.5577 - val\_loss: 1.6446

Epoch 42/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5577 - val\_loss: 1.6446

Epoch 43/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5212  
- loss: 1.7725 - val\_accuracy: 0.5623 - val\_loss: 1.6192

Epoch 44/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5623 - val\_loss: 1.6192

Epoch 45/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.5298  
- loss: 1.7505 - val\_accuracy: 0.5641 - val\_loss: 1.6081

Epoch 46/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5641 - val\_loss: 1.6081

Epoch 47/600  
1000/1000 ————— 40s 40ms/step - accuracy: 0.5330  
- loss: 1.7099 - val\_accuracy: 0.5649 - val\_loss: 1.6139

Epoch 48/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5649 - val\_loss: 1.6139

Epoch 49/600  
1000/1000 ————— 39s 38ms/step - accuracy: 0.5453  
- loss: 1.6780 - val\_accuracy: 0.5797 - val\_loss: 1.5340

Epoch 50/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5797 - val\_loss: 1.5340

Epoch 51/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5478  
- loss: 1.6535 - val\_accuracy: 0.5839 - val\_loss: 1.5338

Epoch 52/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5839 - val\_loss: 1.5338

Epoch 53/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5571  
- loss: 1.6143 - val\_accuracy: 0.5809 - val\_loss: 1.5384

Epoch 54/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5809 - val\_loss: 1.5384

Epoch 55/600  
1000/1000 ————— 39s 38ms/step - accuracy: 0.5654  
- loss: 1.5860 - val\_accuracy: 0.5860 - val\_loss: 1.5412

Epoch 56/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5860 - val\_loss: 1.5412

Epoch 57/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5674  
- loss: 1.5796 - val\_accuracy: 0.5839 - val\_loss: 1.5253

Epoch 58/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5839 - val\_loss: 1.5253

Epoch 59/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5752  
- loss: 1.5414 - val\_accuracy: 0.5929 - val\_loss: 1.4925

Epoch 60/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5929 - val\_loss: 1.4925

Epoch 61/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5773  
- loss: 1.5241 - val\_accuracy: 0.6041 - val\_loss: 1.4637

Epoch 62/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6041 - val\_loss: 1.4637

Epoch 63/600  
1000/1000 ————— 39s 38ms/step - accuracy: 0.5866  
- loss: 1.4970 - val\_accuracy: 0.5987 - val\_loss: 1.4755

Epoch 64/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.5987 - val\_loss: 1.4755

Epoch 65/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5881  
- loss: 1.4767 - val\_accuracy: 0.6031 - val\_loss: 1.4516

Epoch 66/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6031 - val\_loss: 1.4516

Epoch 67/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.6023  
- loss: 1.4336 - val\_accuracy: 0.6063 - val\_loss: 1.4343

Epoch 68/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6063 - val\_loss: 1.4343

Epoch 69/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.5995  
- loss: 1.4272 - val\_accuracy: 0.6047 - val\_loss: 1.4483

Epoch 70/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6047 - val\_loss: 1.4483

Epoch 71/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.6048  
- loss: 1.4037 - val\_accuracy: 0.6025 - val\_loss: 1.4570

Epoch 72/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6025 - val\_loss: 1.4570

Epoch 73/600  
1000/1000 ————— 39s 38ms/step - accuracy: 0.6090  
- loss: 1.3815 - val\_accuracy: 0.6078 - val\_loss: 1.4405

Epoch 74/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6078 - val\_loss: 1.4405

Epoch 75/600  
1000/1000 ————— 40s 40ms/step - accuracy: 0.6211  
- loss: 1.3490 - val\_accuracy: 0.6145 - val\_loss: 1.4152

Epoch 76/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6145 - val\_loss: 1.4152

Epoch 77/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.6208  
- loss: 1.3319 - val\_accuracy: 0.6135 - val\_loss: 1.4037

Epoch 78/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6135 - val\_loss: 1.4037

Epoch 79/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.6316  
- loss: 1.2991 - val\_accuracy: 0.6193 - val\_loss: 1.3880

Epoch 80/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6193 - val\_loss: 1.3880

Epoch 81/600  
1000/1000 ————— 40s 39ms/step - accuracy: 0.6367  
- loss: 1.2768 - val\_accuracy: 0.6153 - val\_loss: 1.3998

Epoch 82/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6153 - val\_loss: 1.3998

Epoch 83/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.6394  
- loss: 1.2687 - val\_accuracy: 0.6176 - val\_loss: 1.4091

Epoch 84/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6176 - val\_loss: 1.4091

Epoch 85/600  
1000/1000 ————— 39s 39ms/step - accuracy: 0.6428  
- loss: 1.2562 - val\_accuracy: 0.6180 - val\_loss: 1.3961

Epoch 86/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6180 - val\_loss: 1.3961

Epoch 87/600  
1000/1000 ————— 38s 38ms/step - accuracy: 0.6432  
- loss: 1.2382 - val\_accuracy: 0.6254 - val\_loss: 1.3634

Epoch 88/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6254 - val\_loss: 1.3634

Epoch 89/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6523  
- loss: 1.2095 - val\_accuracy: 0.6185 - val\_loss: 1.3929

Epoch 90/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6185 - val\_loss: 1.3929

Epoch 91/600  
1000/1000 ————— 38s 38ms/step - accuracy: 0.6576  
- loss: 1.1940 - val\_accuracy: 0.6284 - val\_loss: 1.3665

Epoch 92/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6284 - val\_loss: 1.3665

Epoch 93/600  
1000/1000 ————— 38s 38ms/step - accuracy: 0.6621  
- loss: 1.1833 - val\_accuracy: 0.6385 - val\_loss: 1.3389

Epoch 94/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6385 - val\_loss: 1.3389

Epoch 95/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6621  
- loss: 1.1579 - val\_accuracy: 0.6387 - val\_loss: 1.3340

Epoch 96/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6387 - val\_loss: 1.3340

Epoch 97/600  
1000/1000 ————— 38s 37ms/step - accuracy: 0.6697  
- loss: 1.1463 - val\_accuracy: 0.6342 - val\_loss: 1.3409

Epoch 98/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6342 - val\_loss: 1.3409

Epoch 99/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6716  
- loss: 1.1252 - val\_accuracy: 0.6381 - val\_loss: 1.3399

Epoch 100/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6381 - val\_loss: 1.3399

Epoch 101/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6771  
- loss: 1.1106 - val\_accuracy: 0.6396 - val\_loss: 1.3241

Epoch 102/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6396 - val\_loss: 1.3241

Epoch 103/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6885  
- loss: 1.0751 - val\_accuracy: 0.6380 - val\_loss: 1.3497

Epoch 104/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6380 - val\_loss: 1.3497

Epoch 105/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6884  
- loss: 1.0709 - val\_accuracy: 0.6444 - val\_loss: 1.3227

Epoch 106/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6444 - val\_loss: 1.3227

Epoch 107/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6863  
- loss: 1.0744 - val\_accuracy: 0.6412 - val\_loss: 1.3215

Epoch 108/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6412 - val\_loss: 1.3215

Epoch 109/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6954  
- loss: 1.0426 - val\_accuracy: 0.6373 - val\_loss: 1.3374

Epoch 110/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6373 - val\_loss: 1.3374

Epoch 111/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6983  
- loss: 1.0315 - val\_accuracy: 0.6413 - val\_loss: 1.3399

Epoch 112/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6413 - val\_loss: 1.3399

Epoch 113/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.6978  
- loss: 1.0335 - val\_accuracy: 0.6417 - val\_loss: 1.3300

Epoch 114/600  
1000/1000 ————— 2s 2ms/step - accuracy:  
0.0000e+00 - loss: 0.0000e+00 - val\_accuracy: 0.6417 - val\_loss: 1.3300

Epoch 115/600  
1000/1000 ————— 37s 37ms/step - accuracy: 0.7018  
- loss: 1.0163 - val\_accuracy: 0.6387 - val\_loss: 1.3867

Epoch 115: early stopping  
Restoring model weights from the end of the best epoch: 107.  
Test loss: 1.3215161561965942 / Test accuracy: 0.6412000060081482

## **Efficientnet - B3:**

```
model = create_model('efficientnet_b3', pretrained=False, num_classes=num_classes)
```

Below are the observation found:

```
Epoch 81/120, Loss: 0.0008397004608532025, accuracy: 0.00%
6721.0s  88  Epoch 82/120, Loss: 0.0008494730037121291, accuracy: 0.00%
6802.5s  89  Epoch 83/120, Loss: 0.0010659214445534221, accuracy: 0.00%
6884.1s  90  Epoch 84/120, Loss: 0.0009784858276488412, accuracy: 0.00%
6965.7s  91  Epoch 85/120, Loss: 0.001248635907393668, accuracy: 0.00%
7047.2s  92  Epoch 86/120, Loss: 0.001824407042995172, accuracy: 0.00%
7128.9s  93  Epoch 87/120, Loss: 0.0021520739340568197, accuracy: 0.00%
7210.4s  94  Epoch 88/120, Loss: 0.0022604803016292863, accuracy: 0.00%
7292.1s  95  Epoch 89/120, Loss: 0.0023133259165672515, accuracy: 0.00%
7373.6s  96  Epoch 90/120, Loss: 0.04077925217703362, accuracy: 0.00%
7455.3s  97  Epoch 91/120, Loss: 0.13161427555959435, accuracy: 0.00%
7536.8s  98  Epoch 92/120, Loss: 0.13119682580043998
7618.4s  99  Epoch 93/120, Loss: 0.10303213900025887
7700.0s  100 Epoch 94/120, Loss: 0.08429403640801393
7781.6s  101 Epoch 95/120, Loss: 0.07916979428050638
7863.2s  102 Epoch 96/120, Loss: 0.07520879628254643
7944.8s  103 Epoch 97/120, Loss: 0.06565497622271127
8026.4s  104 Epoch 98/120, Loss: 0.05923773659558236
8108.0s  105 Epoch 99/120, Loss: 0.12180037037292614
8189.8s  106 Epoch 100/120, Loss: 0.11119969401367102
```